

DWAIT BHATT

 dwaitbhatt.com  linkedin.com/in/dwait-bhatt  dwait.bhatt@gmail.com

SUMMARY

PhD student in machine learning and robotics at UCSD with 3 years of AI R&D experience at Samsung Research and ISRO. Focused on making real-world reinforcement learning feasible on physical robots by improving sample-efficiency and reliability.

SKILLS

Relevant Coursework: ML for Robotics, Sensing & Estimation in Robotics, Deep Generative Models, Advanced Computer Vision, Probabilistic Reasoning, Search & Optimization, Information Theory, Linear Systems, Recommender Systems

Programming Languages: Python, Matlab, C++

Libraries: PyTorch, ManiSkill, Robosuite, OpenCV, Open3D, xArm API

EDUCATION

University of California, San Diego

Ph.D. in Electrical & Computer Engineering - Machine Learning & Data Science

Expected Graduation: June 2029

San Diego, California

University of California, San Diego

M.S. in Electrical & Computer Engineering - Intelligent Systems, Robotics & Control

June 2024

Birla Institute of Technology and Science, Pilani (BITS Pilani)

B.E. (Hons.) in Electrical & Electronics Engineering

June 2019

Pilani, Rajasthan

EXPERIENCE

Existential Robotics Laboratory

Graduate Researcher

January 2023 - Present

San Diego, California

- Conducted a large-scale empirical study of 100+ demonstration-augmented RL algorithms, yielding novel hybrid methods and a simple MCQ pretraining recipe for sample-efficient online RL finetuning. [Paper](#) under review at ICRA.
- Developed a latent policy learning framework to enable manipulation skill transfer across diverse robot embodiments; co-authored [paper](#) submitted to RA-L.
- Created a sim-to-real pipeline from scratch for executing robot policies learned in simulation on a physical robot arm. Successfully executed tasks like stacking and pick-and-place on an xArm6 robot with sub-5mm error.

Samsung R&D Institute India

Senior Machine Learning Engineer, On-Device AI

July 2019 - August 2022

Bangalore, Karnataka

- Spearheaded the development of a [patented](#) application launch time prediction model, from a python POC to an on-device implementation in C++, which was used to boost CPU for an optimal duration to make application launches 6% faster while extending battery life by 2 hours.
- Authored a [research paper](#) on this work and presented it at IEEE IJCNN 2022 held in Padua, Italy.
- Automated app launch video analysis using a novel algorithm based on structural similarity between consecutive video frames, saving over 40 man-days annually and significantly enhancing operational efficiency.
- Mentored a team of four undergraduate students through a sequence prediction project for Samsung's industry-academia collaboration program, PRISM.

Machine Learning Engineer, On-Device AI

- Designed a reinforcement learning agent that trained on-device to control an Android smartphone's kernel tunables, reducing the decoding time of encrypted models by 10% on Samsung's flagship phones.

Indian Space Research Organisation (ISRO)

Research Trainee

September 2021 - January 2022

Ahmedabad, Gujarat (Remote)

- Implemented and evaluated a policy-gradient-based deep reinforcement learning approach for adaptive image patch selection, preserving 90% CIFAR-10 accuracy under patch dropping. This implementation was subsequently used in downstream work on satellite and drone imagery for efficient poverty estimation.

PROJECTS

Headset-free Robot Teleoperation | Python

December 2023 - January 2024

- Independently designed and [implemented](#) a VR teleoperation system for real-time control of an xArm6 robot.

Curiosity-driven Robot Reinforcement Learning | PyTorch

January 2023 – March 2023

- Introduced curiosity-based intrinsic motivation in a PPO robot policy via Random Network Distillation. Found that naively adding such exploration hurt performance for robotic tasks in contrast to RND's success in playing video games.

Multimodal Material Estimation | PyTorch

April 2024 – June 2024

- Predicted object material from images and audio of an object being struck by a drumstick by aligning audio and image embeddings to llama-3 token embeddings using cross-attention. The multimodal approach achieved 20% higher accuracy than using only images, proving that audio cues are crucial for object material estimation.

Musical Accompaniment Generation using Language Modeling | PyTorch

April 2023 – June 2023

- Designed a novel architecture for accompaniment generation conditioned on a vocal track using language modeling.
- Achieved semantically meaningful generation (FAD 4.07) with a budget of under \$50.